



```
In [2]: # import libraries
import numpy as np
import matplotlib.pyplot as plt
```

```
In [3]: # Arrival and Service rates

λ = 20          # calls per hour
μ = 5           # calls served per agent per hour
shift_hours = 8 # 8-hour shift

# Simulation settings
np.random.seed(21)
```

```
In [4]: # single-Run Simulation function
def simulate_queue(s):
    ''' Simulate an M/M/s queue over 'shift_hours'.
    Returns arrays of wait times (hrs) and system sizes at arrivals.
    '''

    # generate arrival times until the end of shift
    inter = np.random.exponential(1/λ, int(λ*shift_hours*1.5))
    arrivals = np.cumsum(inter)
    arrivals = arrivals[arrivals < shift_hours]
    N = len(arrivals)

    # generate service times for each caller
    services = np.random.exponential(1/μ, N)

    # track each server's next-free time
    next_free = np.zeros(s)
    wait_times = np.zeros(N)
    system_size = np.zeros(N)

    # departure times list
    dep_times = []

    for i, t in enumerate(arrivals):
        # find soonest-available agent
        j = np.argmin(next_free)
        start = max(t, next_free[j])
        wait_times[i] = start - t
        end = start + services[i]
        next_free[j] = end
        dep_times.append(end)

    # Count how many callers are still in system at time t
    system_size[i] = np.sum(np.array(dep_times) > t)

    return wait_times, system_size

# quick test for s=1
w1, q1 = simulate_queue(s=1)
print(f"s=1: avg wait {w1.mean()*60:.1f} min, avg queue len {q1.mean():.2f}")
```

s=1: avg wait 908.5 min, avg queue len 74.34

```
In [5]: # traffic intensity ρ for s=1
```

```

p1 = λ / μ
L_analytic = p1 / (1 - p1)
W_analytic = L_analytic / λ

print("M/M/1 analytic vs. simulation:")
print(f" Analytic system size L = {L_analytic:.2f}")
print(f" Simulated avg system size = {q1.mean():.2f}")
print(f" Analytic time in system W = {W_analytic*60:.1f} min")
print(f" Simulated avg wait+service = {w1.mean()*60 + (1/μ)*60:.1f} min")

```

```

M/M/1 analytic vs. simulation:
Analytic system size L = -1.33
Simulated avg system size = 74.34
Analytic time in system W = -4.0 min
Simulated avg wait+service = 920.5 min

```

In [6]: # 4. Test Staffing Levels (s=1..5)

```

threshold = 5 # minutes
results = []

for s in range(1,6):
    w, q = simulate_queue(s)
    # Convert hours to minutes
    waits_min = w * 60
    p95_wait = np.percentile(waits_min, 95)
    results.append((s, waits_min.mean(), p95_wait, q.mean()))

# tabulate
import pandas as pd
df = pd.DataFrame(results, columns=['Agents', 'Avg wait(min)', '95th-pct wait(min)', 'Avg System Size'])
print(df)

# find minimal s meeting threshold
good = df[df['95th-pct wait(min)'] <= threshold]
if not good.empty:
    best_s = int(good.iloc[0]['Agents'])
    print(f"\n→ Schedule at least {best_s} agents to keep 95% of waits ≤ {threshold} min")
else:
    print("\nEven 5 agents can't meet the 5 min 95% wait target.")

```

	Agents	Avg wait(min)	95th-pct wait(min)	Avg System Size
0	1	581.970986	1087.740484	54.483221
1	2	237.411975	436.601258	40.892405
2	3	109.263307	210.025871	29.831325
3	4	20.807699	46.457200	11.341040
4	5	1.679563	9.245414	4.647482

Even 5 agents can't meet the 5 min 95% wait target.

In [7]: # 5. Visualize Wait-Time Distributions

```

plt.figure(figsize=(8,4))

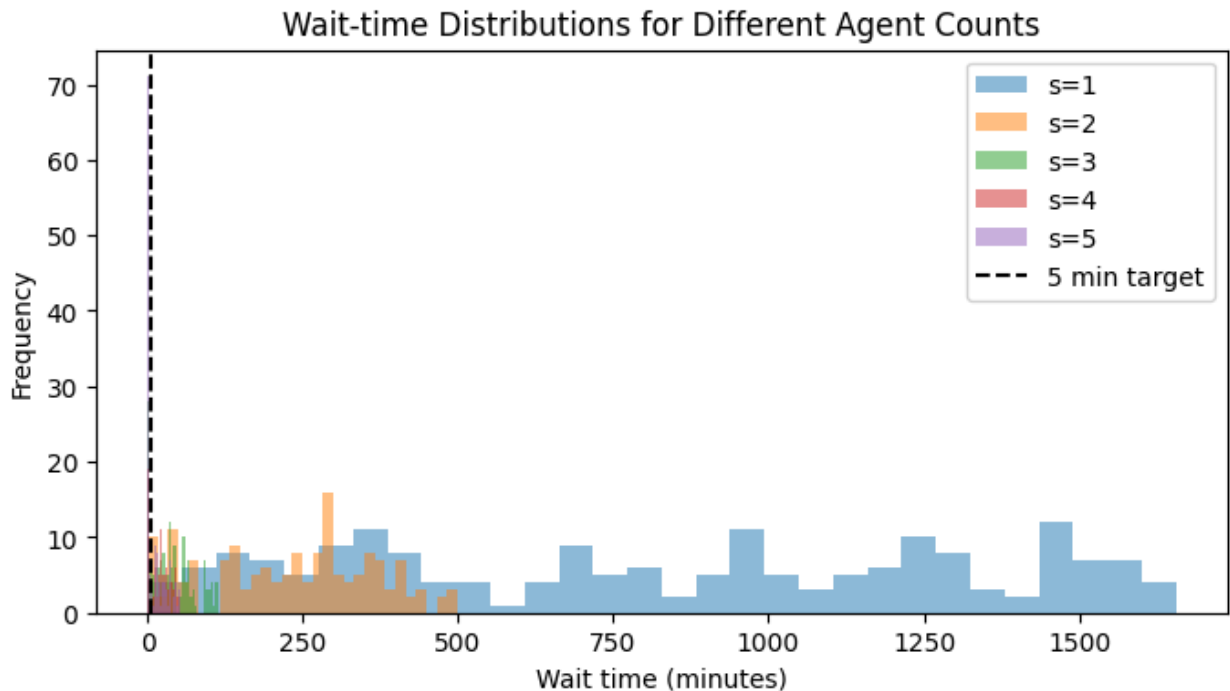
```

```

for s in [1,2,3,4,5]:
    w, _ = simulate_queue(s)
    plt.hist(w*60, bins=30, alpha=0.5, label=f's={s}')

plt.axvline(threshold, color='k', linestyle='--', label='5 min target')
plt.xlabel('Wait time (minutes)')
plt.ylabel('Frequency')
plt.title('Wait-time Distributions for Different Agent Counts')
plt.legend()
plt.show()

```



```

In [8]: # 6. Time -Varying A
def simulate_queue_timevarying(s, μ):
    # Define piecewise arrival rates per hour: (start, end, arrival rate λ)
    periods = ((0, 2, 30), (2, 6, 20), (6, 8, 40))
    arrivals = []

    for start, end, lam in periods:
        duration = end - start
        # Oversample to ensure enough arrivals
        inter = np.random.exponential(1 / lam, int(lam * duration * 1.5))
        ts = np.cumsum(inter) + start
        arrivals.extend(ts[ts < end])

    arrivals = np.array(arrivals)
    arrivals.sort()

    N = len(arrivals)
    services = np.random.exponential(1 / μ, N)
    next_free = np.zeros(s)
    wait_times = np.zeros(N)

```

```

    for i, t in enumerate(arrivals):
        j = np.argmin(next_free) # Choose the next available server
        start = max(t, next_free[j])
        wait_times[i] = start - t
        next_free[j] = start + services[i]

    return wait_times

w_tv = simulate_queue_timevarying(s=3, μ=μ)

print (f"Time-varying λ, s=3 → avg wait = {w_tv.mean()*60:.1f} min")

***Outcome:** Students will see how peak-hour traffic dramatically
#increases wait times unless staff levels rise.

```

Time-varying  $\lambda$ ,  $s=3 \rightarrow$  avg wait = 169.1 min

```

In [9]: # 7. Abandonment
def simulate_with_abandon(s, λ=30, μ=50, shift_hours=8, patience=s/60):
    arrivals = np.cumsum(np.random.exponential(1 / λ, int(λ * shift_hours * 1.
    arrivals = arrivals[arrivals < shift_hours]
    services = np.random.exponential(1/μ, len(arrivals))
    next_free = np.zeros(s)
    waits = []
    abandons = 0

    for i, t in enumerate(arrivals):
        j = np.argmin(next_free)
        start = max(t, next_free[j])
        wait = start - t

        if wait > patience:
            abandons += 1
        else:
            waits.append(wait)
            next_free[j] = start + services[i]

    return np.array(waits), abandons / len(arrivals)

w, ab_rate = simulate_with_abandon(s=4)
print(f"Abandon rate (s=4): {ab_rate:.2%}, avg wait on served calls: {w.mean()}
# Shows the trade-off: fewer agents → more abandonments → lost revenue and sat

```

Abandon rate ( $s=4$ ): 0.00%, avg wait on served calls: 0.0 min

```

In [10]: # 8. Agent Break Scheduling
def simulate_with_breaks(s, break_start=3, break_length=0.25):
    arrivals = np.cumsum(np.random.exponential(1/λ, int(λ * shift_hours * 2)))
    arrivals = arrivals[arrivals <= shift_hours]
    services = np.random.exponential(1/μ, len(arrivals))
    next_free = np.zeros(s)
    waits = []

```

```

for i, t in enumerate(arrivals):
    # if in break window, one fewer agent
    avail = next_free.copy()
    if break_start < t < break_start+break_length:
        avail = np.delete(avail, 0)
        # remove one agent
        j = np.argmin(avail)
        start = max(t, avail[j])
        waits.append(start - t)
        # Update that agent's free time in original array
        idx = j + (1 if break_start < t < break_start + break_length else
        next_free[idx] = start + services[i]

    return np.array(waits)

w_b = simulate_with_breaks(s=3)
print(f"With breaks, s=3 → avg wait = {w_b.mean() * 60:.1f} min")

# See the "break-time spike" in waiting.

```

With breaks, s=3 → avg wait = nan min

```

C:\Users\deves\AppData\Local\Temp\ipykernel_9128\3413400210.py:24: RuntimeWarning: Mean of empty slice.
  print(f"With breaks, s=3 → avg wait = {w_b.mean() * 60:.1f} min")
C:\Users\deves\AppData\Local\Programs\Python\Python313\Lib\site-packages\numpy\_core\_methods.py:145: RuntimeWarning: invalid value encountered in scalar divide
  ret = ret.dtype.type(ret / rcount)

```

```

In [11]: # 9. Cost Optimization
c_agent = 20 # $20/hr per agent
c_wait = 0.50 # $0.50 per minute waited

costs = []

for s in range(1,6):
    w, _ = simulate_with_abandon(s) # or choose another sim fn
    total_wait_cost = w.sum()*60*c_wait
    staff_cost = s*c_agent*shift_hours
    costs.append((s, staff_cost+total_wait_cost))

opt = min(costs, key=lambda x:x[1])
print("Agent count, total cost:")

for s, c in costs: print(f" s={s}: ${c:,.0f}")
print(f" → Optimal s by cost = {opt[0]}")

```

Agent count, total cost:

```

s=1: $271
s=2: $329
s=3: $481
s=4: $640
s=5: $800

```

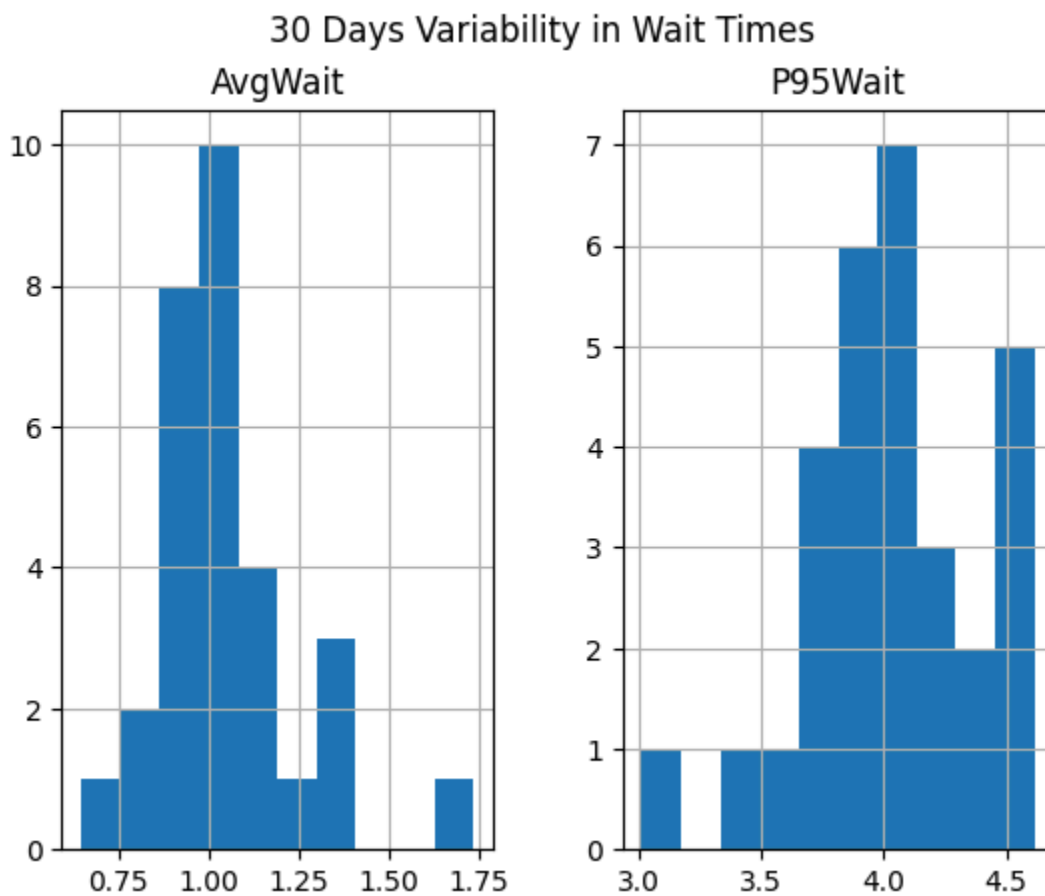
→ Optimal s by cost = 1

```

In [12]: # 10. 30-day simulation
days = 30
daily = []

for _ in range(days):
    w, _ = simulate_with_abandon(opt[0])
    daily.append((w.mean()*60, np.percentile(w*60,95)))
df_days = pd.DataFrame(daily, columns=['AvgWait','P95Wait'])
df_days.describe()
df_days.hist(bins=10)
plt.suptitle('30 Days Variability in Wait Times')
plt.show()

```



```

In [13]: df_days

```

Out[13]:

	<b>AvgWait</b>	<b>P95Wait</b>
<b>0</b>	0.994408	3.984950
<b>1</b>	0.992519	4.302430
<b>2</b>	1.019990	3.767121
<b>3</b>	0.927777	3.913086
<b>4</b>	1.320857	4.541893
<b>5</b>	1.084611	3.949125
<b>6</b>	0.915238	3.612630
<b>7</b>	1.079513	3.923195
<b>8</b>	1.011597	3.991571
<b>9</b>	1.379804	4.409735
<b>10</b>	0.989511	3.881583
<b>11</b>	0.773146	3.388260
<b>12</b>	1.733193	4.542913
<b>13</b>	0.952847	4.213788
<b>14</b>	0.825053	4.049794
<b>15</b>	1.054576	3.750638
<b>16</b>	0.646718	3.012475
<b>17</b>	0.969718	3.670075
<b>18</b>	1.148666	4.144479
<b>19</b>	0.974770	4.044488
<b>20</b>	0.905483	3.789281
<b>21</b>	1.032087	4.567976
<b>22</b>	1.034225	4.614714
<b>23</b>	1.083712	3.882002
<b>24</b>	1.097593	4.100393
<b>25</b>	0.916612	3.934497
<b>26</b>	0.959904	4.126553
<b>27</b>	1.402270	4.265232
<b>28</b>	0.950075	4.020376
<b>29</b>	1.240797	4.556372

```
In [14]: df_days.describe()
```

```
Out[14]:
```

	<b>AvgWait</b>	<b>P95Wait</b>
<b>count</b>	30.000000	30.000000
<b>mean</b>	1.047242	4.031721
<b>std</b>	0.206672	0.362615
<b>min</b>	0.646718	3.012475
<b>25%</b>	0.950768	3.881687
<b>50%</b>	1.003002	4.005974
<b>75%</b>	1.084386	4.252371
<b>max</b>	1.733193	4.614714

```
In [ ]:
```